

# CKA认证练习.note

创建时间: 2020/01/17 09:00

更新时间: 2020/01/20 10:38

来源: <https://www.jxhs.me/2019/11/17/CKA%E9%A2%98%E5%BA%93/>

---

## 第一题 日志 (5%)

**Monitor the logs of Pod foobar** and Extract log lines corresponding to **error unable-to-access-website** Write them to /opt/KULM00201/foobar

监控Pod名称为foobar的日志, 并提取错误unable-to-access-website对应的日志行, 将它们写入/opt/KULM00201/foobar

```
[root@master ~]# kubectl logs foobar | grep unable-to-access-website > /opt/KULM00201/foobar
```

## 第二题 排序 (3%)

**List all PVs sorted by name**, saving the full kubectl output to /opt/kUCC0010/my\_volumes. Use kubectl 's own functionality for sorting the output, and do not mainpulate it any further

列出环境内所有的pv并以name字段排序 (使用 kubectl 自带排序功能)

```
[root@master ~]# kubectl get pv --sort-by=.metadata.name > /opt/kUCC0010/my_volumes
```

**扩展:** 若要按capacity排序列出所有persistent volumes, 将完整的kubectl输出保存到/opt/kUCC0010/my\_volumes(使用 kubectl 的自带功能对输出排序, 不要做其他任何处理)

```
[root@master ~]# kubectl get pv -o yaml #查看pv的--sort-by可自定义对象
[root@master ~]# kubectl get pv --sort-by=.spec.capacity.storage > /opt/kUCC0010/my_volumes
```

## 第三题 DaemonSet编排 (3%)

Ensure a single instance of Pod nginx is running on each node of the kubernetes cluster where nginx also represents the image name which has to be used.

**Do no override any taints currently in place.**

Use Daemonset to complete this task and use ds.kusc00201 as Daemonset name.

确保在kubernetes集群的每个节点上都运行Pod nginx的单个实例，其中nginx还表示要使用的映像名称。切勿覆盖当前存在的任何taint。

使用 DaemonSet 完成此任务，并使用 ds.kusc00201 作为 DaemonSet 名称。

```
[root@master ~]# kubectl run nginx --image=nginx --dry-run -o yaml >
```

```
./ds.kusc00201.yaml
```

```
[root@master ~]# vi ./ds.kusc00201.yaml
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  creationTimestamp: null
```

```
  labels:
```

```
    run: ds.kusc00201
```

```
  name: ds.kusc00201
```

```
spec:
```

```
# replicas: 1
```

```
  selector:
```

```
    matchLabels:
```

```
      run: nginx
```

```
# strategy: {}
```

```
  template:
```

```
    metadata:
```

```
      creationTimestamp: null
```

```
      labels:
```

```
        run: nginx
```

```
    spec:
```

```
      containers:
```

```
        - image: nginx
```

```
          name: nginx
```

```
# resources: {}
```

```
#status: {}
```

```
[root@master ~]# kubectl apply -f ./ds.kusc00201.yaml
```

#### 第四题 init容器 (7%)

Perform the following tasks **Add an init container to lumpy-koala** (which has been defined in spec file `/opt/kucc00100/pod-sepc-KUCC00100.yaml`)

**The init container should create an empty file named `/workdir/calm.txt`.**

**If `/workdir/calm.txt` is not be detected, the Pod should exit.** Once the spec file has been updated with the init container definition, the Pod should be created.

向lumpy-koala添加一个init容器(已在`/opt/kucc00100/pod-spec-kucc00100.yaml`中定义)

init容器应该创建一个名为`/workdir/calm.txt`的空文件

如果`/workdir/calm.txt`没有检测到, Pod应该退出, 一旦使用init容器定义更新了spec文件, 就创建Pod

```
[root@master ~]# vi /opt/kucc00100/pod-spec-kucc00100.yaml
```

*containers.image*加入如下内容 (保持对其) :

```
command: ["/bin/sh", "-c", "if [ -f /workdir/calm.txt ]; then sleep 100000; else exit 1; fi"]
```

*initContainers.image*加入如下内容 (保持对其) :

```
command: ["/bin/sh", "-c", "touch /workdir/calm.txt"]
```

```
[root@master ~]# kubectl apply -f /opt/kucc00100/pod-spec-kucc00100.yaml
```

## 第五题 多个容器的pod的创建 (7%)

Create a pod named **kucc4** with a single container for each of the following images running inside (there may be between 1 and 4 images specified):**nginx + redis + memcached + consul**

创建一个名为kucc4的pod, 1个Pod容器其中包含4个镜像: nginx+redis+memcached+consul

```
[root@master ~]# kubectl run kucc4 --image=nginx --dry-run -o yaml > kucc4.yaml
```

```
[root@master ~]# vi ./kucc4.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  creationTimestamp: null
```

```
  labels:
```

```
    run: kucc4
```

```
  name: kucc4
```

```
#spec:
```

```
# replicas: 1
```

```
# selector:
```

```

# matchLabels:
#   run: kucc4
# strategy: {}
# template:
#   metadata:
#     creationTimestamp: null
#     labels:
#       run: kucc4
spec:
  containers:
    - image: nginx
      name: nginx
    - image: redis
      name: redis
    - image: memcached
      name: memcached
    - image: consul
      name: consul
#   resources: {}
#status: {}

```

[root@master ~]# kubectl apply -f kucc4.yaml

注意：如果指定要用busybox镜像一定要指定sleep，不然一直重启

```

spec:
  containers:
    - name: busybox
      image: busybox:1.28
      command:
        - sleep
        - "3600"

```

## 第六题 pod的调度 (2%)

Schedule a Pod as follows:

Name:nginx-kusc00101

Image:nginx

NodeSelector:disk=ssd

创建一个pod名称为nginx，并将其调度到标签节点为 disk=ssd上

```
[root@master ~]# kubectl get node --show-labels | grep ssd #查看ssd的标签归属在哪个node
```

```
[root@master ~]# vi node2.yaml #打开官网链接复制模板
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-kusc00101
```

```
  labels:
```

```
    env: test
```

```
spec:
```

```
  containers:
```

```
    - name: nginx
```

```
      image: nginx
```

```
      imagePullPolicy: IfNotPresent
```

```
  nodeSelector:
```

```
    disktype: ssd
```

```
[root@master ~]# kubectl apply -f ./node2.yaml
```

注意：如需要自己手动打便签，执行如下

```
[root@master ~]# kubectl label nodes node2 disktype=ssd
```

## 第七题 Deployment 资源的更新 (4%)

Create a deployment as follows

Name: nginx-app

Using container nginx with version 1.11.9-alpine

The deployment should contain 3 replicas

Next, **deploy the app with new version 1.12.0-alpine by performing a rolling update and record that update.**

Finally, rollback that update to the previous version 1.11.9-alpine

按照以下方式创建部署

名称:nginx-app

使用1.11.9-alpine版本的容器nginx

deployment应该包含3个副本

接下来，**使用新版本1.12.0-alpine部署应用程序，执行滚动更新并记录更新。**

最后，将更新回滚到前一个版本1.11.9-alpine

```
[root@master ~]# kubectl run nginx-app --image=nginx:1.11.9-alpine --replicas=3
[root@master ~]# kubectl set image deployment nginx-app nginx-app=nginx:1.12.0-alpine --record=true
[root@master ~]# kubectl rollout undo deployment/nginx-app
[root@master ~]# kubectl rollout status -w deployment nginx-app
[root@master ~]# kubectl rollout history deployment/nginx-app
```

注意，查看详情可以执行

```
[root@master ~]# kubectl get deployment nginx-app -o yaml
```

## 第八题 Service (4%)

Create and configure the service front-end-service

so it 's accessible through NodePort/ClusterIP and routes to the existing pod named front-end

使用front-end-service服务，将名为front-end的pod，用NodePort/ClusterIP的方式发布出来

```
[root@master ~]# kubectl expose pod front-end --name=front-end-service --
type=NodePort --port=80
[root@master ~]# kubectl get svc
```

## 第九题 新分区创建Pod (3%)

Create a Pod as follows:

Name: jenkins

Using image: jenkins

In a new Kubernetes namespce named website-frontend

按如下要求创建一个 Pod:

名称: jenkins

使用 image: jenkins

在名为 website-frontend 的 新 kubernetes namespace 中

```
[root@master ~]# kubectl get ns
[root@master ~]# kubectl create ns website-frontend
```

```
[root@master ~]# kubectl run jenkins --image=jenkins --namespace=website-frontend --generator=run-pod/v1
```

## 第十题 Deployment创建 (3%)

Create a deployment spce file that will:

launch 7 replicas of the redis image with the label:app\_env\_stage=dev

Deployment name: kua100201

Save a copy of this spec file to /opt/KUAL00201/deploy\_spec.yaml

When you are done,clean up (delete) any new k8s API objects that you produced during this task

创建一个deployment文件，文件将：

启动7个redis镜像副本，镜像标签是：app\_env\_stage=dev

deployment名称：kual00201

将规范文件的副本保存到/opt/KUAL002001/deploy\_spc.yamle (or .json)

完成后，清理(删除)在此任务期间生成的任何新的k8s API对象

```
[root@master ~]# mkdir /opt/KUAL00201
```

```
[root@master ~]# kubectl run kua100201 --image=redis --replicas=7 --
```

```
labels=app_env_stage=dev --dry-run -o yaml > /opt/KUAL00201/deploy_spec.yaml
```

```
[root@master ~]# kubectl apply -f /opt/KUAL00201/deploy_spec.yaml
```

```
[root@master ~]# kubectl delete -f /opt/KUAL00201/deploy_spec.yaml
```

## 第十一题 统计Service中的pod (3%)

Create a file /opt/KUCC00302/kucc00302.txt that lists all pods that implement Service foo in Namespce production

The format of the file should be one pod name per line

创建一个文件/opt/KUCC00302/kucc00302.txt,

其中列出在Namespce 为Production中实现Service 为foo的所有pod

文件的格式应该是每行一个pod名称

```
[root@master ~]# mkdir /opt/KUCC00302
```

```
[root@master ~]# touch /opt/KUCC00302/kucc00302.txt
```

```
[root@master ~]# kubectl get svc --show-labels -n production | grep foo > /opt/KUCC00302/kucc00302.txt #手动删除其他多余内容，只需要保留第一列的Pod名称即可
```

## 第十二题 secret (9%)

Create a kubernetes Secret as follows:

Name: super-secret

Credential: alice or username:bob

Create a Pod named **pod-secrets-via-file** using the redis image which **mounts a secret named super-secret at /secrets**

Create a second Pod named **pod-secrets-via-env** using the redis image, **which exports credential/username as TOPSECRET/CREDENTIALS**

创建一个secret，使用以下：

名字：super-secret

Credential：alice or username：bob

创建一个pod名为pod-secrets-via-file 使用redis镜像，挂载名为super-secret的挂载路径/secrets

使用redis镜像创建第二个Pod名称Pod-secrets-via-env，使用credential/username 的方式，对应的变量为：TOPSECRET/CREDENTIALS

```
[root@master ~]# kubectl create secret generic super-secret --from-literal=credential=alice --from-literal=username=bob
```

```
[root@master ~]# kubectl get secret
```

```
[root@master ~]# vi super-secret.yaml #打开官网，复制两个模板env修改
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod-secrets-via-file
```

```
  labels:
```

```
    name: pod-secrets-via-file
```

```
spec:
```

```
  volumes:
```

```
    - name: super-secret
```

```
      secret:
```

```
        secretName: super-secret
```

```
  containers:
```

```
    - name: pod-secrets-via-file
```

```
      image: redis
```



*volumeMounts:*

- ***name: super-secret***

*readOnly: true*

***mountPath: /secrets***

*apiVersion: v1*

*kind: Pod*

*metadata:*

***name: pod-secrets-via-env***

*spec:*

*containers:*

- ***name: pod-secrets-via-env***

***image: redis***

*env:*

- ***name: TOPSECRET***

*valueFrom:*

*secretKeyRef:*

***name: super-secret***

***key: credential***

- ***name: CREDENTIALS***

*valueFrom:*

*secretKeyRef:*

***name: super-secret***

***key: username***

***# restartPolicy: Never***

[root@master ~]# kubectl apply -f ./super-secret.yaml

[root@master ~]# kubectl describe pod-secrets-via-file

[root@master ~]# kubectl describe pod-secrets-via-env

### 第十三题 EmptyDir (4%)

按如下要求创建一个 Pod:

名称: non-persistent-redis

Container image: redis

Persistent volume name: cache-control

Mount Path: /data/redis

应在 staging namespace 中发布, 且该 volume 必须不能是永久的。

[root@master ~]# kubectl get ns

[root@master ~]# kubectl create ns staging

[root@master ~]# vi non-persistent-redis.yaml #打开官网连接直接复制

```
apiVersion: v1
kind: Pod
metadata:
  name: non-persistent-redis
  namespace: staging
spec:
  containers:
    - image: redis
      name: non-persistent-redis
  volumeMounts:
    - mountPath: /data/redis
      name: cache-control
  volumes:
    - name: cache-control
      emptyDir: {}
```

```
[root@master ~]# kubectl apply -f non-persistent-redis.yaml
```

```
[root@master ~]# kubectl describe pod -n staging non-persistent-redis
```

#### 第十四题 Scale: 扩缩容 (1%)

Scale the deployment webserver to 6 pods

将部署webserver扩展到6个pod

```
[root@master ~]# kubectl get deployment
```

```
[root@master ~]# kubectl scale deployment/webserver --replicas=6
```

```
[root@master ~]# kubectl get deployment
```

#### 第十五题 统计node是ready状态 (2%)

check to see how many nodes are ready (not including nodes tainted NoSchedule) and write the number to /opt/nodenum

检查有多少nodes是ready状态, (不包含node的污点, 没有调度的), 写入数量到 /opt/nodenum

```
[root@master ~]# kubectl get nodes | grep Ready #注意有多少行
```

```
[root@master ~]# kubectl describe node | grep Taints #注意NoSchedule的有多少行  
[root@master ~]# echo 行数 > /opt/nodenum #Ready总行-NoSchedule的行=最终有效行
```

## 第十六题 统计pod的CPU (2%)

From the pod label name=cpu-utilizer, find **pods running high CPU** workloads and write the name of the pod consuming most CPU to the file /opt/cpu.txt

从标签为name=cpu-utilizer的所有pod里面，找出cpu使用最高的那个pod，并写入到/opt/cpu.txt（这个文件已经存在）

```
[root@master ~]# kubectl get pod -A --show-labels | grep cpu-utilizer  
[root@master ~]# kubectl top pod -n spacename podname  
[root@master ~]# echo podname >> /opt/cpu.txt #只需要保留pod的名称即可
```

## 第十七题 nslookup查找service和pod的DNS记录 (7%)

Create a deployment as follows

Name: nginx-dns

Exposed via a service: nginx-dns

**Ensure that the service & pod are accessible via their respective DNS records**

The container(s) within any pod(s) running as a part of this deployment should use the nginx image

Next, use the utility nslookup to look up the DNS records of the service & pod and write the output to **/opt/service.dns** and /opt/pod.dns respectively

**Ensure you use the busybox:1.28 image (or earlier) for any testing**, as the latest release has an upstream bug which impacts the use of nslookup

创建一个deployment 名为nginx-dns，发布一个服务名为: nginx-dns

确保service 和pod可以通过各自的DNS记录访问

作为deployment的一部分运行的任何pod中的容器都应该使用nginx镜像

接下来，使用实用程序nslookup查找service和pod的DNS记录，并将输出分别写入/opt/service.dns和/opt/pod.dns

确保您使用busybox:1.28 image(或更早版本)进行任何测试，最新版本有一个upstream bug，影响nslookup的使用。

```
[root@master ~]# kubectl run nginx-dns --image=nginx
```

```
[root@master ~]# kubectl expose deployment nginx-dns --port=80
[root@master ~]# kubectl run busybox -it --image=busybox:1.28 sh
[root@master ~]# kubectl exec -ti busyboxname -- nslookup nginx-dns > /opt/service.dns
[root@master ~]# kubectl exec -ti busyboxname -- nslookup <Pod ip> > /opt/pod.dns
```

## 第十八题 etcd快照 (7%)

Create a snapshot of **the etcd instance running at http://127.0.0.1:2379** saving the **snapshot to the file path /data/backup/etcd-snapshot.db**

**The etcd instance is running etcd version 3.2.18**

The following TLS certificates/key are supplied for connecting to the server with etcdctl

**CA certificate: /opt/KUCM00302/ca.crt**

**Client certificate: /opt/KUCM00302/etcd-client.crt**

**Client key: /opt/KUCM00302/etcd-client.key**

```
[root@master ~]# etcdctl --endpoints=http://127.0.0.1:2379 --ca-
file=/opt/KUCM00302/ca.crt --certfile=/opt/KUCM00302/etcd-client.crt --
key=/opt/KUCM00302/etcd-client.key snapshot save /data/backup/etcd-snapshot.db
```

如有需要指定将输出结果保存到某一个文件

```
[root@master ~]# etcdctl snapshot status /etcd-snapshot.db > XXX.txt
```

## 第十九题 使用 kubectl drain 从集群中移除节点 (4%)

Set configuration context \$ kubectl config use-context ek8s

Set the node labelled with name=ek8s-node-1 as **unavailable** and **reschedule all the pods running on it**

把标签为name=ek8s-node-1的node设置为unavailable和重新安排所有运行在上面的pods

```
[root@master ~]# kubectl get nodes --show-labels -A | grep name=ek8s-node-1
[root@master ~]# kubectl drain nodename
```

注意：若要恢复节点，可以执行如下

```
[root@master ~]# kubectl uncordon
```

如果直接drain会出错，需要添加--ignore-daemonsets --delete-local-data参数

```
[root@master ~]#kubectl drain node node1 --ignore-daemonsets --delete-local-data
```

## 第二十题 修复notready状态的节点 (4%)

A Kubernetes worker node, labelled with name=wk8s-node-0 is in state NotReady. Investigate why this is the case, and perform any appropriate steps to bring the node to a Ready state, ensuring that any changes are made permanent.

Hints:

You can ssh to the failed node using `$ ssh wk8s-node-0`

You can assume elevated privileges on the node with the following command `$ sudo -i`

wk8s集群里面有一个标签为wk8s-node-0的节点是notready状态，找到原因，恢复Ready状态，所做的改变要是持久的

```
[root@master ~]# kubectl get node #查看哪个node为notready状态，ssh连接上去
```

```
[root@master ~]# sudo -i
```

```
[root@master ~]# systemctl status kubelet
```

```
[root@master ~]# systemctl start kubelet
```

```
[root@master ~]# systemctl enable kubelet
```

## 第二十一题 Create static Pods (4%)

configure **the kubelet systemd managed service**, on the node labelled with **name=wk8s-node-1**, to launch a pod containing a single container of image **nginx** named **myservice** automatically.

**Any spec file** required should be placed in the **/etc/kubernetes/manifests** directory on the node

Hints:

You can ssh to the failed node using `$ ssh wk8s-node-0`

You can assume elevated privileges on the node with the following command `$ sudo -i`

配置一个kubelet 系统管理的服务，在标签为name1=wk8s-node-1的节点上配置，要包含一个POD名为myservice的镜像nginx 容器。

所需的任何特定文件应放在/etc/kubernetes/manifests 的节点文件夹内

```
[root@node1 ~]# ssh wk8s-node-0
```

```
[root@node1 ~]# sudo -i
```

```
[root@node1 ~]# kubectl run myservice --image=nginx --generator=run-pod/v1 --dry-run -o yaml > /etc/kubernetes/manifests/wk8s-node-0.yaml
[root@node1 ~]# vi /var/lib/kubelet/config.yaml #查看staticPodPath:
/etc/kubernetes/manifests
[root@node1 ~]# systemctl daemon-reload
[root@node1 ~]# systemctl restart kubelet
[root@node1 ~]# systemctl enable kubelet
```

@@@@第二十二题 给出一个集群，将节点node1添加到集群中（8%）

第二十三题 集群故障排查（kubelet配置的静态Pod路径）（4%）

Given a partially-functioning Kubernetes cluster, identify symptoms of failure on the cluster.

**Determine the node, the failing service** and take actions to bring up the failed service and restore the health of the cluster.

**Ensure that any changes are made permanently.**

The worker node in this cluster is labelled with name=bk8s-node-0

Hints:

You can ssh to the relevant nodes using \$ **ssh \$(NODE)** where **\$(NODE)** is one of **bk8s-master-0** or **bk8s-node-0**

You can assume elevated privileges on any node in the cluster with the following command  
\$ **sudo -i**

给定一个部分功能正常的Kubernetes集群，识别集群上的故障症状。

确定节点、failing服务器并采取启动失败的服务并恢复集群的健康状态，确保永久地进行任何更改。

这个集群中的工作节点被标记为name=bk8 -node-0

情况一：

```
[root@master ~]# kubectl get node,cs #查看是否是群集组件API问题，一般报6443.....
[root@master ~]# cd /etc/kubernetes/manifests/ #确定群集是否用kubeadm方式部署的，存在文件
[root@master ~]# vi /var/lib/kubelet/config.yaml #确定staticPodPath路径是否正确
[root@master ~]# systemctl restart kubelet
[root@master ~]# kubectl get node,cs
```

情况二：

```
[root@master ~]# kubectl get node,cs
```

```
[root@master ~]# ssh master
```

```
[root@master ~]# systemctl restart kube-manager-controller.service
```

## 第二十四题 持久卷 (3%)

Create a persistent volume with name app-config of capacity 1Gi and access mode ReadWriteOnce.

The type of volume is hostPath and its location is /srv/app-config

创建一个持久卷，名称为app-config，容量为1Gi，访问模式为ReadWriteOnce  
卷的类型是hostPath，它的位置是/srv/app-config

```
[root@node1 ~]# vi persistent.yaml
```

```
apiVersion: v1
```

```
kind: PersistentVolume
```

```
metadata:
```

```
  name: app-config
```

```
spec:
```

```
  capacity:
```

```
    storage: 1Gi
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
  hostPath:
```

```
    path: /srv/app-config
```

```
[root@node1 ~]# kubectl apply -f persistent.yaml
```